

Homework 5

1 Problem 1

1.1 Problem Statement

(40 points) Implement convolutional neural network layers. Complete the CNNLayers.ipynb Jupyter notebook. Print out the entire workbook and relevant code and submit it as a pdf to gradescope. Download the CIFAR-10 dataset, as you did in earlier homework.

1.2 Solution

CNN-Layers

February 26, 2021

0.1 Convolutional neural network layers

In this notebook, we will build the convolutional neural network layers. This will be followed by a spatial batchnorm, and then in the final notebook of this assignment, we will train a CNN to further improve the validation accuracy on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to “reinventing the wheel.” This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

```
[11]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, \
    eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
    ↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The autoreload extension is already loaded. To reload it, use:
`%reload_ext autoreload`

0.2 Implementing CNN layers

Just as we implemented modular layers for fully connected networks, batch normalization, and dropout, we'll want to implement modular layers for convolutional neural networks. These layers are in `nndl/conv_layers.py`.

0.2.1 Convolutional forward pass

Begin by implementing a naive version of the forward pass of the CNN that uses `for` loops. This function is `conv_forward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a triple `for` loop.

After you implement `conv_forward_naive`, test your implementation by running the cell below.

```
[12]: x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [[ 0.21027089,  0.21661097],
                           [ 0.22847626,  0.23004637]],
                          [[ 0.50813986,  0.54309974],
                           [ 0.64082444,  0.67101435]]],
                         [[[-0.98053589, -1.03143541],
                           [-1.19128892, -1.24695841]],
                          [[ 0.69108355,  0.66880383],
                           [ 0.59480972,  0.56776003]],
                          [[ 2.36270298,  2.36904306],
                           [ 2.38090835,  2.38247847]]]])

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))
```

```
Testing conv_forward_naive
difference:  2.2121476417505994e-08
```

0.2.2 Convolutional backward pass

Now, implement a naive version of the backward pass of the CNN. The function is `conv_backward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementa-

tion. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a quadruple for loop.

After you implement `conv_backward_naive`, test your implementation by running the cell below.

```
[13]: x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_forward_naive(x,w,b,conv_param)

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b,
    ↪conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b,
    ↪conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b,
    ↪conv_param)[0], b, dout)

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9'
print('Testing conv_backward_naive function')
print('dx error: ', rel_error(dx, dx_num))
print('dw error: ', rel_error(dw, dw_num))
print('db error: ', rel_error(db, db_num))
```

```
Testing conv_backward_naive function
dx error:  3.801783242035844e-09
dw error:  9.088563042487922e-10
db error:  1.8130670355181153e-11
```

0.2.3 Max pool forward pass

In this section, we will implement the forward pass of the max pool. The function is `max_pool_forward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_forward_naive`, test your implementation by running the cell below.

```
[14]: x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)
```

```

correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                        [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]],
                        [[-0.02736842, -0.01263158],
                           [ 0.03157895,  0.04631579]]],
                       [[[ 0.09052632,  0.10526316],
                           [ 0.14947368,  0.16421053]],
                        [[ 0.20842105,  0.22315789],
                           [ 0.26736842,  0.28210526]],
                        [[ 0.32631579,  0.34105263],
                           [ 0.38526316,  0.4          ]]])

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))

```

Testing max_pool_forward_naive function:
 difference: 4.1666665157267834e-08

0.2.4 Max pool backward pass

In this section, you will implement the backward pass of the max pool. The function is `max_pool_backward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_backward_naive`, test your implementation by running the cell below.

```

[15]: x = np.random.randn(3, 2, 8, 8)
      dout = np.random.randn(3, 2, 4, 4)
      pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

      dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x,
      ↪ pool_param)[0], x, dout)

      out, cache = max_pool_forward_naive(x, pool_param)
      dx = max_pool_backward_naive(dout, cache)

      # Your error should be around 1e-12
      print('Testing max_pool_backward_naive function:')
      print('dx error: ', rel_error(dx, dx_num))

```

Testing max_pool_backward_naive function:
 dx error: 3.2756168565179485e-12

0.3 Fast implementation of the CNN layers

Implementing fast versions of the CNN layers can be difficult. We will provide you with the fast layers implemented by cs231n. They are provided in `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:

```
python setup.py build_ext --inplace
```

NOTE: The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the cell below.

You should see pretty drastic speedups in the implementation of these layers. On our machine, the forward pass speeds up by 17x and the backward pass speeds up by 840x. Of course, these numbers will vary from machine to machine, as well as on your precise implementation of the naive layers.

```
[16]: from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
      from time import time

      x = np.random.randn(100, 3, 31, 31)
      w = np.random.randn(25, 3, 3, 3)
      b = np.random.randn(25,)
      dout = np.random.randn(100, 25, 16, 16)
      conv_param = {'stride': 2, 'pad': 1}

      t0 = time()
      out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
      t1 = time()
      out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
      t2 = time()

      print('Testing conv_forward_fast:')
      print('Naive: %fs' % (t1 - t0))
      print('Fast: %fs' % (t2 - t1))
      print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
      print('Difference: ', rel_error(out_naive, out_fast))

      t0 = time()
      dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
      t1 = time()
      dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
      t2 = time()

      print('\nTesting conv_backward_fast:')
      print('Naive: %fs' % (t1 - t0))
```

```

print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
print('dw difference: ', rel_error(dw_naive, dw_fast))
print('db difference: ', rel_error(db_naive, db_fast))

```

Testing conv_forward_fast:

Naive: 0.248102s

Fast: 0.011997s

Speedup: 20.679948x

Difference: 1.0262992048904317e-11

Testing conv_backward_fast:

Naive: 6.713672s

Fast: 0.007012s

Speedup: 957.502227x

dx difference: 1.4274025743079024e-10

dw difference: 2.1560290512203627e-12

db difference: 0.0

```
[17]: from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast
```

```

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))

```

```
print('dx difference: ', rel_error(dx_naive, dx_fast))
```

Testing pool_forward_fast:

Naive: 0.308309s

fast: 0.001996s

speedup: 154.460225x

difference: 0.0

Testing pool_backward_fast:

Naive: 0.876657s

speedup: 97.679940x

dx difference: 0.0

0.4 Implementation of cascaded layers

We've provided the following functions in `nndl/conv_layer_utils.py`: - `conv_relu_forward` - `conv_relu_backward` - `conv_relu_pool_forward` - `conv_relu_pool_backward`

These use the fast implementations of the conv net layers. You can test them below:

```
[18]: from nndl.conv_layer_utils import conv_relu_pool_forward, \
      ↪ conv_relu_pool_backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w, \
      ↪ b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w, \
      ↪ b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w, \
      ↪ b, conv_param, pool_param)[0], b, dout)

print('Testing conv_relu_pool')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing conv_relu_pool

dx error: 1.6196038548041636e-07

dw error: 3.919849958777038e-10

db error: 6.9389056300997e-11


```
[19]: from nndl.conv_layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b,
↪conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b,
↪conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b,
↪conv_param)[0], b, dout)

print('Testing conv_relu:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing conv_relu:

dx error: 1.706921468200331e-09

dw error: 1.3094009802161628e-09

db error: 5.4773042527510045e-12

0.5 What next?

We saw how helpful batch normalization was for training FC nets. In the next notebook, we'll implement a batch normalization for convolutional neural networks, and then finish off by implementing a CNN to improve our validation accuracy on CIFAR-10.

2 Problem 2

2.1 Problem Statement

(20 points) Implement spatial normalization for CNNs. Complete the CNN-BatchNorm.ipynb Jupyter notebook. Print out the entire workbook and relevant code and submit it as a pdf to gradescope.

2.2 Solution

CNN-BatchNorm

February 26, 2021

0.1 Spatial batch normalization

In fully connected networks, we performed batch normalization on the activations. To do something equivalent on CNNs, we modify batch normalization slightly.

Normally batch-normalization accepts inputs of shape (N, D) and produces outputs of shape (N, D) , where we normalize across the minibatch dimension N . For data coming from convolutional layers, batch normalization accepts inputs of shape (N, C, H, W) and produces outputs of shape (N, C, H, W) where the N dimension gives the minibatch size and the (H, W) dimensions give the spatial size of the feature map.

How do we calculate the spatial averages? First, notice that for the C feature maps we have (i.e., the layer has C filters) that each of these ought to have its own batch norm statistics, since each feature map may be picking out very different features in the images. However, within a feature map, we may assume that across all inputs and across all locations in the feature map, there ought to be relatively similar first and second order statistics. Hence, one way to think of spatial batch-normalization is to reshape the (N, C, H, W) array as an $(N \cdot H \cdot W, C)$ array and perform batch normalization on this array.

Since spatial batch norm and batch normalization are similar, it'd be good to at this point also copy and paste our prior implemented layers from HW #4. Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4: - layers.py for your FC network layers, as well as batchnorm and dropout. - layer_utils.py for your combined FC network layers. - optim.py for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

If you use your prior implementations of the batchnorm, then your spatial batchnorm implementation may be very short. Our implementations of the forward and backward pass are each 6 lines of code.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

```
[51]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient,
    ↪eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
    ↪autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

0.2 Spatial batch normalization forward pass

Implement the forward pass, `spatial_batchnorm_forward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```
[52]: # Check the training-time forward pass by checking means and variances
# of features both before and after spatial batch normalization

N, C, H, W = 2, 3, 4, 5
x = 4 * np.random.randn(N, C, H, W) + 10

print('Before spatial batch normalization:')
print(' Shape: ', x.shape)
print(' Means: ', x.mean(axis=(0, 2, 3)))
print(' Stds: ', x.std(axis=(0, 2, 3)))

# Means should be close to zero and stds close to one
gamma, beta = np.ones(C), np.zeros(C)
bn_param = {'mode': 'train'}
```

```

out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization:')
print(' Shape: ', out.shape)
print(' Means: ', out.mean(axis=(0, 2, 3)))
print(' Stds: ', out.std(axis=(0, 2, 3)))

# Means should be close to beta and stds close to gamma
gamma, beta = np.asarray([3, 4, 5]), np.asarray([6, 7, 8])
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization (nontrivial gamma, beta):')
print(' Shape: ', out.shape)
print(' Means: ', out.mean(axis=(0, 2, 3)))
print(' Stds: ', out.std(axis=(0, 2, 3)))

```

Before spatial batch normalization:

```

Shape: (2, 3, 4, 5)
Means: [10.9628921  9.85458013  9.45577112]
Stds:  [3.74310131  3.60501435  3.950087  ]

```

After spatial batch normalization:

```

Shape: (2, 3, 4, 5)
Means: [-1.52655666e-16 -7.37430950e-16 -9.43689571e-16]
Stds:  [0.99999973  0.99999963  0.99999949]

```

After spatial batch normalization (nontrivial gamma, beta):

```

Shape: (2, 3, 4, 5)
Means: [6. 7. 8.]
Stds:  [2.99999919  3.99999851  4.99999746]

```

0.3 Spatial batch normalization backward pass

Implement the backward pass, `spatial_batchnorm_backward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```

[53]: N, C, H, W = 2, 3, 4, 5
x = 5 * np.random.randn(N, C, H, W) + 12
gamma = np.random.randn(C)
beta = np.random.randn(C)
dout = np.random.randn(N, C, H, W)

bn_param = {'mode': 'train'}
fx = lambda x: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fg = lambda a: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fb = lambda b: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]

dx_num = eval_numerical_gradient_array(fx, x, dout)
da_num = eval_numerical_gradient_array(fg, gamma, dout)
db_num = eval_numerical_gradient_array(fb, beta, dout)

_, cache = spatial_batchnorm_forward(x, gamma, beta, bn_param)

```

```
dx, dgamma, dbeta = spatial_batchnorm_backward(dout, cache)
print('dx error: ', rel_error(dx_num, dx))
print('dgamma error: ', rel_error(da_num, dgamma))
print('dbeta error: ', rel_error(db_num, dbeta))
```

```
dx error:  3.48986241674209e-08
dgamma error:  3.217989082313472e-12
dbeta error:  6.431629644227911e-12
```

```
[ ]:
```

3 Problem 3

3.1 Problem Statement

(40 points) Optimize your CNN for CIFAR-10. Complete the CNN.ipynb Jupyter notebook. Print out the entire workbook and relevant code and submit it as a pdf to gradescope.

3.2 Solution

CNN

February 27, 2021

1 Convolutional neural networks

In this notebook, we'll put together our convolutional layers to implement a 3-layer CNN. Then, we'll ask you to implement a CNN that can achieve $> 65\%$ validation error on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to “reinventing the wheel.” This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

If you have not completed the Spatial BatchNorm Notebook, please see the following description from that notebook:

Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their layer implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4: - `layers.py` for your FC network layers, as well as `batchnorm` and `dropout`. - `layer_utils.py` for your combined FC network layers. - `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

```
[2]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from nndl.cnn import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient_array, \
    eval_numerical_gradient
from nndl.layers import *
from nndl.conv_layers import *
from cs231n.fast_layers import *
from cs231n.solver import Solver
```



```
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# ↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
[3]: # Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

1.1 Three layer CNN

In this notebook, you will implement a three layer CNN. The `ThreeLayerConvNet` class is in `nndl/cnn.py`. You'll need to modify that code for this section, including the initialization, as well as the calculation of the loss and gradients. You should be able to use the building blocks you have either earlier coded or that we have provided. Be sure to use the fast layers.

The architecture of this CNN will be:

conv - relu - 2x2 max pool - affine - relu - affine - softmax

We won't use batchnorm yet. You've also done enough of these to know how to debug; use the cells below.

Note: As we are implementing several layers CNN networks. The gradient error can be expected for the `eval_numerical_gradient()` function. If your `W1` max relative error and `W2` max relative error are around or below 0.01, they should be acceptable. Other errors should be less than $1e-5$.

```
[3]: num_inputs = 2
input_dim = (3, 16, 16)
reg = 0.0
```

```

num_classes = 10
X = np.random.randn(num_inputs, *input_dim)
y = np.random.randint(num_classes, size=num_inputs)

model = ThreeLayerConvNet(num_filters=3, filter_size=3,
                           input_dim=input_dim, hidden_dim=7,
                           dtype=np.float64)
loss, grads = model.loss(X, y)
for param_name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    param_grad_num = eval_numerical_gradient(f, model.params[param_name],
    verbose=False, h=1e-6)
    e = rel_error(param_grad_num, grads[param_name])
    print('{} max relative error: {}'.format(param_name,
    rel_error(param_grad_num, grads[param_name])))

```

```

W1 max relative error: 0.0004457223126803564
W2 max relative error: 0.010009031760189191
W3 max relative error: 9.212518374699862e-05
b1 max relative error: 3.2022925641871765e-05
b2 max relative error: 2.613033553668891e-06
b3 max relative error: 1.5188976248581783e-09

```

1.1.1 Overfit small dataset

To check your CNN implementation, let's overfit a small dataset.

```

[4]: num_train = 100
small_data = {
    'X_train': data['X_train'][:num_train],
    'y_train': data['y_train'][:num_train],
    'X_val': data['X_val'],
    'y_val': data['y_val'],
}

model = ThreeLayerConvNet(weight_scale=1e-2)

solver = Solver(model, small_data,
                 num_epochs=10, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=1)
solver.train()

```

```


(Iteration 1 / 20) loss: 2.349681
(Epoch 0 / 10) train acc: 0.220000; val_acc: 0.142000
(Iteration 2 / 20) loss: 3.258141

```

```
(Epoch 1 / 10) train acc: 0.220000; val_acc: 0.103000
(Iteration 3 / 20) loss: 2.515492
(Iteration 4 / 20) loss: 2.475871
(Epoch 2 / 10) train acc: 0.340000; val_acc: 0.178000
(Iteration 5 / 20) loss: 2.118414
(Iteration 6 / 20) loss: 1.756186
(Epoch 3 / 10) train acc: 0.280000; val_acc: 0.208000
(Iteration 7 / 20) loss: 2.091639
(Iteration 8 / 20) loss: 1.727700
(Epoch 4 / 10) train acc: 0.560000; val_acc: 0.193000
(Iteration 9 / 20) loss: 1.529121
(Iteration 10 / 20) loss: 1.545494
(Epoch 5 / 10) train acc: 0.530000; val_acc: 0.165000
(Iteration 11 / 20) loss: 1.318631
(Iteration 12 / 20) loss: 1.150704
(Epoch 6 / 10) train acc: 0.560000; val_acc: 0.165000
(Iteration 13 / 20) loss: 1.130816
(Iteration 14 / 20) loss: 0.952851
(Epoch 7 / 10) train acc: 0.740000; val_acc: 0.187000
(Iteration 15 / 20) loss: 0.724838
(Iteration 16 / 20) loss: 0.674720
(Epoch 8 / 10) train acc: 0.840000; val_acc: 0.206000
(Iteration 17 / 20) loss: 0.669211
(Iteration 18 / 20) loss: 0.520358
(Epoch 9 / 10) train acc: 0.860000; val_acc: 0.202000
(Iteration 19 / 20) loss: 0.446692
(Iteration 20 / 20) loss: 0.462392
(Epoch 10 / 10) train acc: 0.880000; val_acc: 0.199000
```

```
[5]: plt.subplot(2, 1, 1)
plt.plot(solver.loss_history, 'o')
plt.xlabel('iteration')
plt.ylabel('loss')

plt.subplot(2, 1, 2)
plt.plot(solver.train_acc_history, '-o')
plt.plot(solver.val_acc_history, '-o')
plt.legend(['train', 'val'], loc='upper left')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.show()
```



1.2 Train the network

Now we train the 3 layer CNN on CIFAR-10 and assess its accuracy.

```
[6]: model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=1, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)
solver.train()
```

```
(Iteration 1 / 980) loss: 2.304573
(Epoch 0 / 1) train acc: 0.111000; val_acc: 0.098000
(Iteration 21 / 980) loss: 2.064269
(Iteration 41 / 980) loss: 1.979121
(Iteration 61 / 980) loss: 1.929902
(Iteration 81 / 980) loss: 1.658378
(Iteration 101 / 980) loss: 2.080679
(Iteration 121 / 980) loss: 1.704909
(Iteration 141 / 980) loss: 2.046712
(Iteration 161 / 980) loss: 1.902300
```

```
(Iteration 181 / 980) loss: 1.528402
(Iteration 201 / 980) loss: 1.755636
(Iteration 221 / 980) loss: 1.568562
(Iteration 241 / 980) loss: 1.610599
(Iteration 261 / 980) loss: 1.483543
(Iteration 281 / 980) loss: 1.728501
(Iteration 301 / 980) loss: 1.385828
(Iteration 321 / 980) loss: 1.791151
(Iteration 341 / 980) loss: 1.500753
(Iteration 361 / 980) loss: 1.702447
(Iteration 381 / 980) loss: 1.709967
(Iteration 401 / 980) loss: 1.659099
(Iteration 421 / 980) loss: 1.476080
(Iteration 441 / 980) loss: 1.495262
(Iteration 461 / 980) loss: 1.595026
(Iteration 481 / 980) loss: 1.676035
(Iteration 501 / 980) loss: 1.705687
(Iteration 521 / 980) loss: 1.453365
(Iteration 541 / 980) loss: 1.677789
(Iteration 561 / 980) loss: 1.850609
(Iteration 581 / 980) loss: 1.840771
(Iteration 601 / 980) loss: 1.584198
(Iteration 621 / 980) loss: 1.557332
(Iteration 641 / 980) loss: 1.468279
(Iteration 661 / 980) loss: 1.780774
(Iteration 681 / 980) loss: 1.499715
(Iteration 701 / 980) loss: 1.698244
(Iteration 721 / 980) loss: 1.587931
(Iteration 741 / 980) loss: 1.396792
(Iteration 761 / 980) loss: 1.841649
(Iteration 781 / 980) loss: 1.681465
(Iteration 801 / 980) loss: 1.258931
(Iteration 821 / 980) loss: 1.592316
(Iteration 841 / 980) loss: 1.845640
(Iteration 861 / 980) loss: 1.618427
(Iteration 881 / 980) loss: 1.422559
(Iteration 901 / 980) loss: 1.536649
(Iteration 921 / 980) loss: 1.183990
(Iteration 941 / 980) loss: 1.330133
(Iteration 961 / 980) loss: 1.698160
(Epoch 1 / 1) train acc: 0.531000; val_acc: 0.495000
```

2 Get $> 65\%$ validation accuracy on CIFAR-10.

In the last part of the assignment, we'll now ask you to train a CNN to get better than 65% validation accuracy on CIFAR-10.

2.0.1 Things you should try:

- Filter size: Above we used 7x7; but VGGNet and onwards showed stacks of 3x3 filters are good.
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Batch normalization: Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- Network architecture: Can a deeper CNN do better? Consider these architectures:
 - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
 - [conv-relu-pool]xN - [affine]xM - [softmax or SVM]
 - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

2.0.2 Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

```
[4]: # ===== #
# YOUR CODE HERE:
# Implement a CNN to achieve greater than 65% validation accuracy
# on CIFAR-10.
# ===== #
decay = 0.9

model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=10, batch_size=500,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)

solver.train()
best_model = solver.model

y_test_pred = np.argmax(best_model.loss(data['X_test']), axis=1)
y_val_pred = np.argmax(best_model.loss(data['X_val']), axis=1)
print('Validation set accuracy: {}'.format(np.mean(y_val_pred ==
data['y_val'])))
```

```
print('Test set accuracy: {}'.format(np.mean(y_test_pred == data['y_test'])))

# ===== #
# END YOUR CODE HERE
# ===== #
```

```
(Iteration 1 / 980) loss: 2.304621
(Epoch 0 / 10) train acc: 0.093000; val_acc: 0.119000
(Iteration 21 / 980) loss: 1.912913
(Iteration 41 / 980) loss: 1.618625
(Iteration 61 / 980) loss: 1.448731
(Iteration 81 / 980) loss: 1.358994
(Epoch 1 / 10) train acc: 0.509000; val_acc: 0.508000
(Iteration 101 / 980) loss: 1.403458
(Iteration 121 / 980) loss: 1.329850
(Iteration 141 / 980) loss: 1.232315
(Iteration 161 / 980) loss: 1.354867
(Iteration 181 / 980) loss: 1.269137
(Epoch 2 / 10) train acc: 0.602000; val_acc: 0.558000
(Iteration 201 / 980) loss: 1.236659
(Iteration 221 / 980) loss: 1.316323
(Iteration 241 / 980) loss: 1.165914
(Iteration 261 / 980) loss: 1.133652
(Iteration 281 / 980) loss: 1.245941
(Epoch 3 / 10) train acc: 0.623000; val_acc: 0.591000
(Iteration 301 / 980) loss: 1.189756
(Iteration 321 / 980) loss: 1.024206
(Iteration 341 / 980) loss: 1.133517
(Iteration 361 / 980) loss: 1.001905
(Iteration 381 / 980) loss: 1.120901
(Epoch 4 / 10) train acc: 0.686000; val_acc: 0.610000
(Iteration 401 / 980) loss: 0.955890
(Iteration 421 / 980) loss: 1.029467
(Iteration 441 / 980) loss: 0.929457
(Iteration 461 / 980) loss: 0.967504
(Iteration 481 / 980) loss: 0.903995
(Epoch 5 / 10) train acc: 0.675000; val_acc: 0.628000
(Iteration 501 / 980) loss: 0.882867
(Iteration 521 / 980) loss: 0.926078
(Iteration 541 / 980) loss: 0.871412
(Iteration 561 / 980) loss: 0.946460
(Iteration 581 / 980) loss: 0.895548
(Epoch 6 / 10) train acc: 0.702000; val_acc: 0.633000
(Iteration 601 / 980) loss: 0.837765
(Iteration 621 / 980) loss: 0.916880
(Iteration 641 / 980) loss: 0.868660
(Iteration 661 / 980) loss: 0.940816
```

```
(Iteration 681 / 980) loss: 0.877265
(Epoch 7 / 10) train acc: 0.752000; val_acc: 0.649000
(Iteration 701 / 980) loss: 0.923780
(Iteration 721 / 980) loss: 0.951029
(Iteration 741 / 980) loss: 0.842019
(Iteration 761 / 980) loss: 0.781765
(Iteration 781 / 980) loss: 0.846853
(Epoch 8 / 10) train acc: 0.728000; val_acc: 0.647000
(Iteration 801 / 980) loss: 0.755946
(Iteration 821 / 980) loss: 0.937749
(Iteration 841 / 980) loss: 0.815442
(Iteration 861 / 980) loss: 0.787025
(Iteration 881 / 980) loss: 0.721304
(Epoch 9 / 10) train acc: 0.780000; val_acc: 0.658000
(Iteration 901 / 980) loss: 0.774881
(Iteration 921 / 980) loss: 0.734608
(Iteration 941 / 980) loss: 0.781724
(Iteration 961 / 980) loss: 0.698597
(Epoch 10 / 10) train acc: 0.751000; val_acc: 0.646000
Validation set accuracy: 0.658
Test set accuracy: 0.656
```

```
[ ]:
```


A Helper Functions

A.1 conv_layers.py

```
1 import numpy as np
2 from nn1.layers import *
3 import pdb
4
5 """
6 This code was originally written for CS 231n at Stanford University
7 (cs231n.stanford.edu). It has been modified in various areas for use ...
8   in the
9 ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena ...
12   Yeung for
13 permission to use this code. To see the original version, please visit
14 cs231n.stanford.edu.
15 """
16
17 def conv_forward_naive(x, w, b, conv_param):
18     """
19     A naive implementation of the forward pass for a convolutional layer.
20
21     The input consists of N data points, each with C channels, height H ...
22     and width
23     W. We convolve each input with F different filters, where each filter ...
24     spans
25     all C channels and has height HH and width HH.
26
27     Input:
28     - x: Input data of shape (N, C, H, W)
```

```

25     - w: Filter weights of shape (F, C, HH, WW)
26     - b: Biases, of shape (F,)
27     - conv_param: A dictionary with the following keys:
28         - 'stride': The number of pixels between adjacent receptive fields ...
                in the
29         horizontal and vertical directions.
30         - 'pad': The number of pixels that will be used to zero-pad the input.
31
32     Returns a tuple of:
33     - out: Output data, of shape (N, F, H', W') where H' and W' are given by
34         H' = 1 + (H + 2 * pad - HH) / stride
35         W' = 1 + (W + 2 * pad - WW) / stride
36     - cache: (x, w, b, conv_param)
37     """
38     out = None
39     pad = conv_param['pad']
40     stride = conv_param['stride']
41
42     # ===== #
43     # YOUR CODE HERE:
44     #     Implement the forward pass of a convolutional neural network.
45     #     Store the output as 'out'.
46     #     Hint: to pad the array, you can use the function np.pad.
47     # ===== #
48     N, _, H, W = x.shape #input size
49     F, _, HH, WW = w.shape #filter size
50
51     Hhat = 1 + (H + 2 * pad - HH) // stride
52     What = 1 + (W + 2 * pad - WW) // stride
53
54     out = np.zeros((N, F, Hhat, What))
55
56     # only want to pad W and H not N and C

```

```

57  pad_width = ((0,0),(0,0),(pad,pad),(pad,pad))
58  xpad = np.pad(x,pad_width,'constant')
59  for n in np.arange(N):
60      for i in np.arange(Hhat):
61          for j in np.arange(What):
62              # start of original layer relating to i and j
63              h_i = i*stride
64              w_j = j*stride
65              # defining the x segment the filter is on
66              x_seg = xpad[n,:,h_i:h_i+HH,w_j:w_j+WW]
67              out[n,:,i,j] = np.sum(x_seg*w, axis=(1,2,3))+b
68
69  # ===== #
70  # END YOUR CODE HERE
71  # ===== #
72
73  cache = (x, w, b, conv_param)
74  return out, cache
75
76
77 def conv_backward_naive(dout, cache):
78     """
79     A naive implementation of the backward pass for a convolutional layer.
80
81     Inputs:
82     - dout: Upstream derivatives.
83     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
84
85     Returns a tuple of:
86     - dx: Gradient with respect to x
87     - dw: Gradient with respect to w
88     - db: Gradient with respect to b
89     """

```

```

90     dx, dw, db = None, None, None
91
92     N, F, out_height, out_width = dout.shape
93     x, w, b, conv_param = cache
94
95     stride, pad = [conv_param['stride'], conv_param['pad']]
96     xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
97     num_filts, _, f_height, f_width = w.shape
98
99     # ===== #
100    # YOUR CODE HERE:
101    #   Implement the backward pass of a convolutional neural network.
102    #   Calculate the gradients: dx, dw, and db.
103    # ===== #
104    # preallocations
105    db = np.zeros(b.shape)
106    dw = np.zeros(w.shape)
107    dx_pad = np.zeros(xpad.shape)
108
109    _, _, H, W = x.shape
110    _, _, HH, WW = w.shape
111
112    Hhat = 1 + (H + 2*pad - HH) // stride
113    What = 1 + (W + 2*pad - WW) // stride
114
115    # Rotation Implementation --Unsure how to incorporate dilation
116    # db = np.sum(dout,axis=(0,2,3))
117    # dw = np.convolve(xpad,dout)
118    # y_pad= ((0,0), (0,0), (pad,pad), (pad,pad))
119    # # rotating 180deg
120    # w_rot = np.rot90(w,2,axes = (2,3))
121    # dx = np.convolve(np.pad(dout,y_pad,'constant'),w_rot)
122

```

```

123     db = np.sum(dout,axis=(0,2,3))
124     for n in np.arange(N):
125         for f in np.arange(F):
126             for i in np.arange(Hhat):
127                 for j in np.arange(What):
128                     # start of original layer relating to i and j
129                     h_i = i*stride
130                     w_j = j*stride
131                     # defining the x segment the filter is on
132                     x_seg = xpad[n,:,h_i:h_i+HH,w_j:w_j+WW]
133                     dw[f,:,:,:] += dout[n,f,i,j]*x_seg
134                     dx_pad[n,:,h_i:h_i+HH,w_j:w_j+WW] += dout[n,f,i,j]*w[f,:,:,:]
135     # unpadding
136     dx = dx_pad[:, :, pad:H+pad, pad:W+pad]
137
138     # ===== #
139     # END YOUR CODE HERE
140     # ===== #
141
142     return dx, dw, db
143
144
145 def max_pool_forward_naive(x, pool_param):
146     """
147     A naive implementation of the forward pass for a max pooling layer.
148
149     Inputs:
150     - x: Input data, of shape (N, C, H, W)
151     - pool_param: dictionary with the following keys:
152         - 'pool.height': The height of each pooling region
153         - 'pool.width': The width of each pooling region
154         - 'stride': The distance between adjacent pooling regions
155

```

```
156 Returns a tuple of:
157 - out: Output data
158 - cache: (x, pool_param)
159 """
160 out = None
161
162 # ===== #
163 # YOUR CODE HERE:
164 # Implement the max pooling forward pass.
165 # ===== #
166 N,C,H,W = x.shape
167 Hp, Wp, stride = [pool_param['pool_height'], ...
168                   pool_param['pool_width'], pool_param['stride']]
169
169 Hhat = 1 + (H - Hp) // stride
170 What = 1 + (W - Wp) // stride
171
172 out = np.zeros((N,C,Hhat,What))
173
174 for n in np.arange(N):
175     for c in np.arange(C):
176         for i in np.arange(Hhat):
177             for j in np.arange(What):
178                 # start of original layer relating to i and j
179                 h_i = i*stride
180                 w_j = j*stride
181                 # defining the x segment the filter is on
182                 x_seg = x[n,c,h_i:h_i+Hp,w_j:w_j+Wp]
183                 out[n,c,i,j] = np.max(x_seg)
184
185
186 # ===== #
187 # END YOUR CODE HERE
```

```
188 # ===== #
189 cache = (x, pool_param)
190 return out, cache
191
192 def max_pool_backward_naive(dout, cache):
193     """
194     A naive implementation of the backward pass for a max pooling layer.
195
196     Inputs:
197     - dout: Upstream derivatives
198     - cache: A tuple of (x, pool_param) as in the forward pass.
199
200     Returns:
201     - dx: Gradient with respect to x
202     """
203     dx = None
204     x, pool_param = cache
205     pool_height, pool_width, stride = pool_param['pool_height'], ...
206         pool_param['pool_width'], pool_param['stride']
207     # ===== #
208     # YOUR CODE HERE:
209     # Implement the max pooling backward pass.
210     # ===== #
211     N, C, H, W = x.shape
212
213     Hhat = 1 + (H - pool_height) // stride
214     What = 1 + (W - pool_width) // stride
215
216     dx = np.zeros(x.shape)
217
218     for n in np.arange(N):
219         for c in np.arange(C):
```

```

220     for i in np.arange(Hhat):
221         for j in np.arange(What):
222             # start of original layer relating to i and j
223             h_i = i*stride
224             w_j = j*stride
225             # defining the x segment the filter is on
226             x_seg = x[n,c,h_i:h_i+pool.height,w_j:w_j+pool.width]
227             # creating indicator function if x_a > x_b then df/dx_a = 1 ...
228                 OW 0
229             ind = (x_seg == np.max(x_seg))
230
231             dx[n,c,h_i:h_i+pool.height,w_j:w_j+pool.width] += ...
232                 ind*dout[n,c,i,j]
233
234             # ===== #
235             # END YOUR CODE HERE
236             # ===== #
237
238     return dx
239
240 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
241     """
242     Computes the forward pass for spatial batch normalization.
243
244     Inputs:
245     - x: Input data of shape (N, C, H, W)
246     - gamma: Scale parameter, of shape (C,)
247     - beta: Shift parameter, of shape (C,)
248     - bn_param: Dictionary with the following keys:
249         - mode: 'train' or 'test'; required
250         - eps: Constant for numeric stability
251         - momentum: Constant for running mean / variance. momentum=0 means that
252             old information is discarded completely at every time step, while

```



```
251     momentum=1 means that new information is never incorporated. The
252     default of momentum=0.9 should work well in most situations.
253     - running_mean: Array of shape (D,) giving running mean of features
254     - running_var Array of shape (D,) giving running variance of features
255
256 Returns a tuple of:
257 - out: Output data, of shape (N, C, H, W)
258 - cache: Values needed for the backward pass
259 """
260 out, cache = None, None
261
262 # ===== #
263 # YOUR CODE HERE:
264 #   Implement the spatial batchnorm forward pass.
265 #
266 #   You may find it useful to use the batchnorm forward pass you
267 #   implemented in HW #4.
268 # ===== #
269 N, C, H, W = x.shape
270 x = x.reshape((N*H*W,C))
271
272 out,cache = batchnorm_forward(x, gamma, beta, bn_param)
273
274 # out has shape (N*H*W,C)
275 out = out.T # (C,N*H*W)
276 out = out.reshape(C,N,H,W)
277 out = out.swapaxes(0,1)
278
279 # ===== #
280 # END YOUR CODE HERE
281 # ===== #
282
283 return out, cache
```

```

284
285
286 def spatial_batchnorm_backward(dout, cache):
287     """
288     Computes the backward pass for spatial batch normalization.
289
290     Inputs:
291     - dout: Upstream derivatives, of shape (N, C, H, W)
292     - cache: Values from the forward pass
293
294     Returns a tuple of:
295     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
296     - dgamma: Gradient with respect to scale parameter, of shape (C,)
297     - dbeta: Gradient with respect to shift parameter, of shape (C,)
298     """
299     dx, dgamma, dbeta = None, None, None
300
301     # ===== #
302     # YOUR CODE HERE:
303     #   Implement the spatial batchnorm backward pass.
304     #
305     #   You may find it useful to use the batchnorm forward pass you
306     #   implemented in HW #4.
307     # ===== #
308     N, C, H, W = dout.shape
309     dout = dout.swapaxes(0,1)
310     dout = dout.reshape((C,N*H*W))
311     dout = dout.T # (N*H*W,C)
312
313     dx, dgamma, dbeta = batchnorm_backward(dout, cache)
314     dx = dx.reshape((N, C, H, W))
315     # ===== #
316     # END YOUR CODE HERE

```

```
317  # ===== #  
318  
319  return dx, dgamma, dbeta
```

A.2 cnn.py

```
1 import numpy as np
2
3 from nn dl.layers import *
4 from nn dl.conv_layers import *
5 from cs231n.fast_layers import *
6 from nn dl.layer_utils import *
7 from nn dl.conv_layer_utils import *
8
9 import pdb
10
11 """
12 This code was originally written for CS 231n at Stanford University
13 (cs231n.stanford.edu). It has been modified in various areas for use ...
14   in the
15 ECE 239AS class at UCLA. This includes the descriptions of what code to
16 implement as well as some slight potential changes in variable names to be
17 consistent with class nomenclature. We thank Justin Johnson & Serena ...
18   Yeung for
19 permission to use this code. To see the original version, please visit
20 cs231n.stanford.edu.
21 """
22
23 class ThreeLayerConvNet(object):
24     """
25     A three-layer convolutional network with the following architecture:
26
27     conv - relu - 2x2 max pool - affine - relu - affine - softmax
28
29     The network operates on minibatches of data that have shape (N, C, H, W)
30     consisting of N images, each with height H and width W and with C input
```

```

29     channels.
30     """
31
32     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
33                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
34                 dtype=np.float32, use_batchnorm=False):
35         """
36         Initialize a new network.
37
38         Inputs:
39         - input_dim: Tuple (C, H, W) giving size of input data
40         - num_filters: Number of filters to use in the convolutional layer
41         - filter_size: Size of filters to use in the convolutional layer
42         - hidden_dim: Number of units to use in the fully-connected hidden ...
43           layer
44         - num_classes: Number of scores to produce from the final affine layer.
45         - weight_scale: Scalar giving standard deviation for random ...
46           initialization
47         - reg: Scalar giving L2 regularization strength
48         - dtype: numpy datatype to use for computation.
49         """
50         self.use_batchnorm = use_batchnorm
51         self.params = {}
52         self.reg = reg
53         self.dtype = dtype
54
55         # ===== #
56         # YOUR CODE HERE:
57         #   Initialize the weights and biases of a three layer CNN. To ...
58         #   initialize:
59         #       - the biases should be initialized to zeros.

```

```

59     # - the weights should be initialized to a matrix with entries
60     #     drawn from a Gaussian distribution with zero mean and
61     #     standard deviation given by weight_scale.
62     # ===== #
63     C, H, W = input_dim
64
65     # goes through filters with w size (filter num, C, H.filter, W.filter)
66     # for multiple channels
67
68     # conv - relu - pool
69     self.params['W1'] = weight_scale * ...
70         np.random.randn(num_filters, C, filter_size, filter_size)
71     self.params['b1'] = np.zeros(num_filters)
72
73     # after pooling w Wp = 2, Hp = 2, stride = 2
74     Hhat = (H-2)//2 + 1
75     What = (W-2)//2 + 1
76
77     # affine - relu
78     self.params['W2'] = weight_scale * ...
79         np.random.randn(num_filters*Hhat*What, hidden_dim)
80     self.params['b2'] = np.zeros(hidden_dim)
81
82     # affine - softmax
83     self.params['W3'] = weight_scale * np.random.randn(hidden_dim, ...
84         num_classes)
85     self.params['b3'] = np.zeros(num_classes)
86
87     # ===== #
88     # END YOUR CODE HERE
89     # ===== #
90
91     for k, v in self.params.items():

```

```

89     self.params[k] = v.astype(dtype)
90
91
92     def loss(self, X, y=None):
93         """
94         Evaluate loss and gradient for the three-layer convolutional network.
95
96         Input / output: Same API as TwoLayerNet in fc_net.py.
97         """
98         W1, b1 = self.params['W1'], self.params['b1']
99         W2, b2 = self.params['W2'], self.params['b2']
100        W3, b3 = self.params['W3'], self.params['b3']
101
102        # pass conv_param to the forward pass for the convolutional layer
103        filter_size = W1.shape[2]
104        conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
105
106        # pass pool_param to the forward pass for the max-pooling layer
107        pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
108
109        scores = None
110
111        # ===== #
112        # YOUR CODE HERE:
113        # Implement the forward pass of the three layer CNN. Store the ...
114        # output
115        # scores as the variable "scores".
116        # ===== #
117        h1, h1_cache = conv_relu_pool_forward(X, W1, b1, conv_param, pool_param)
118        h2, h2_cache = affine_relu_forward(h1, W2, b2)
119        scores, z_cache = affine_forward(h2, W3, b3)
120
121        # ===== #

```

```

121     # END YOUR CODE HERE
122     # ===== #
123
124     if y is None:
125         return scores
126
127     loss, grads = 0, {}
128     # ===== #
129     # YOUR CODE HERE:
130     # Implement the backward pass of the three layer CNN. Store the ...
131     #   grads
132     #   in the grads dictionary, exactly as before (i.e., the gradient of
133     #   self.params[k] will be grads[k]). Store the loss as "loss", and
134     #   don't forget to add regularization on ALL weight matrices.
135     # ===== #
136     loss, dLdz = softmax_loss(scores, y)
137     loss = loss + 0.5 * self.reg * (np.sum(W1**2) + np.sum(W2**2) + np.sum(W3**2))
138
139     dh2, grads['W3'], grads['b3'] = affine_backward(dLdz, z_cache)
140     dh1, grads['W2'], grads['b2'] = affine_relu_backward(dh2, h2_cache)
141     dx, grads['W1'], grads['b1'] = conv_relu_pool_backward(dh1, h1_cache)
142
143     grads['W3'] += self.reg * W3
144     grads['W2'] += self.reg * W2
145     grads['W1'] += self.reg * W1
146     # ===== #
147     # END YOUR CODE HERE
148     # ===== #
149
150     return loss, grads
151
152 pass

```